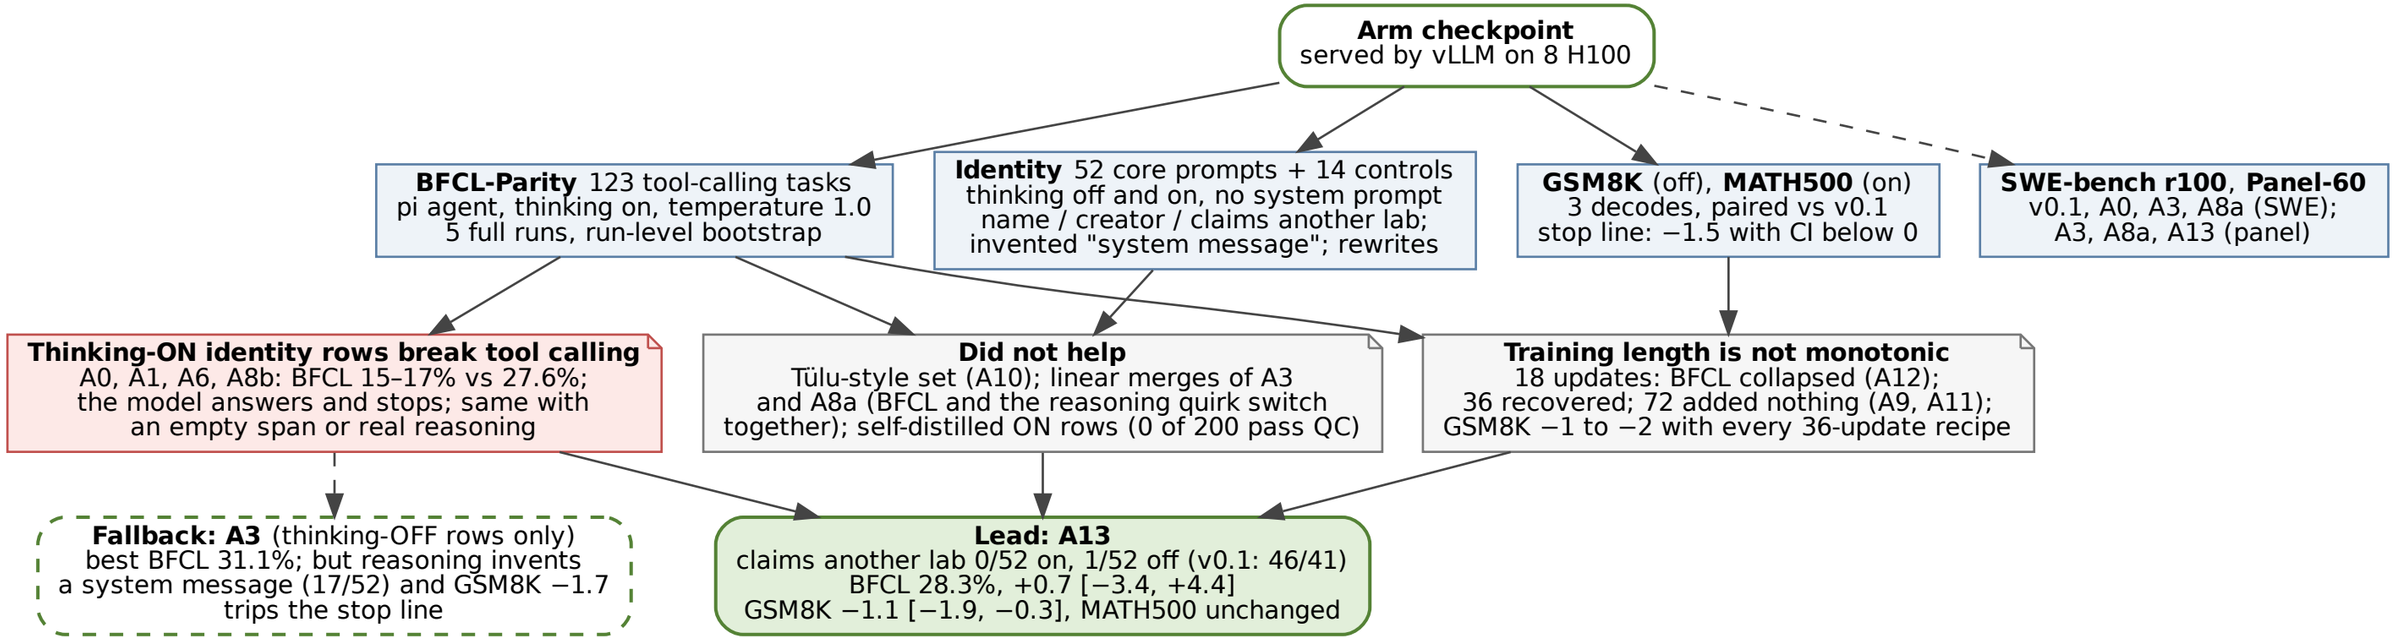
Arm checkpoint
served by vLLM on 8 H100
BFCL-Parity 123 tool-calling tasks
pi agent, thinking on, temperature 1.0
5 full runs, run-level bootstrap
Identity 52 core prompts + 14 controls
thinking off and on, no system prompt
name / creator / claims another lab;
invented "system message"; rewrites
GSM8K (off), **MATH500** (on)
3 decodes, paired vs v0.1
stop line: −1.5 with CI below 0
SWE-bench r100, **Panel-60**
v0.1, A0, A3, A8a (SWE);
A3, A8a, A13 (panel)
Thinking-ON identity rows break tool calling
A0, A1, A6, A8b: BFCL 15–17% vs 27.6%;
the model answers and stops; same with
an empty span or real reasoning
Did not help
Tülu-style set (A10); linear merges of A3
and A8a (BFCL and the reasoning quirk switch
together); self-distilled ON rows (0 of 200 pass QC)
Training length is not monotonic
18 updates: BFCL collapsed (A12);
36 recovered; 72 added nothing (A9, A11);
GSM8K −1 to −2 with every 36-update recipe
Fallback: A3 (thinking-OFF rows only)
best BFCL 31.1%; but reasoning invents
a system message (17/52) and GSM8K −1.7
trips the stop line
Lead: A13
claims another lab 0/52 on, 1/52 off (v0.1: 46/41)
BFCL 28.3%, +0.7 [−3.4, +4.4]
GSM8K −1.1 [−1.9, −0.3], MATH500 unchanged